

```
In [1]: import boto3
import pandas as pd
import numpy as np
import os
from datetime import datetime
import requests
import zipfile
import io
import re
import time
import urllib.parse
```

```
In [2]: s3 = boto3.resource(
    service_name = 's3',
    region_name = 'ap-south-1',
    aws_access_key_id='AKIAQ-----HT5W5',
    aws_secret_access_key = 'upWVXX6----ucuZu72IH----78CeMwb---jxka'
)
```

```
In [3]: for bucket in s3.buckets.all():
    print(bucket.name)
```

```
portfolio-return-data
portfolio-return-raw-bucket
```

## 1. Data Ingestion :

- a. Download raw data from source.
- b. Extract
- c. Parse the files into structured format

```
In [4]: def get_zip_links():
    years = range(2017, 2026)
    base = "https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives"

    links = []
    for y in years:
        folder = f"08 {y} Update"
        for ext in ["TXT"]:
```

```

        url = f"{base}/{folder}/ftp/6_Portfolios_2x3_{ext}.zip"
        links.append(url)

    return links

zip_links = get_zip_links()
zip_links

```

Out[4]:

```

['https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2017 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2018 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2019 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2020 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2021 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2022 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2023 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2024 Update/ftp/6_Portfolios_2x3_TXT.zip',
 'https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2025 Update/ftp/6_Portfolios_2x3_TXT.zip']

```

In [5]:

```

output_dir = r"C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv"
os.makedirs(output_dir, exist_ok=True)

markers = [
    "Average Value Weighted Returns -- Monthly",
    "Average Equal Weighted Returns -- Monthly",
    "Average Value Weighted Returns -- Annual"
]

def parse_fama_french_text(text):
    lines = text.splitlines()
    positions = {}

    for i, line in enumerate(lines):
        for m in markers:

```

```

        if m in line:
            positions[m] = i

    tables = {}
    marker_list = sorted(positions.items(), key=lambda x: x[1])

    for idx in range(len(marker_list)-1):
        name, start = marker_list[idx]
        _, end = marker_list[idx + 1]

        block = lines[start + 2:end]
        rows = []

        for line in block:
            parts = line.strip().split()
            if parts and parts[0].isdigit():
                rows.append(parts)

        if rows:
            df = pd.DataFrame(rows, columns=[
                "Date",
                "Small_Low", "Small_2", "Small_High",
                "Big_Low", "Big_2", "Big_High"
            ])

            for col in df.columns[1:]:
                df[col] = pd.to_numeric(df[col], errors="coerce")

            df["Date"] = pd.to_datetime(df["Date"], format="%Y%m", errors="coerce")
            tables[name] = df

    return tables

# MAIN LOOP
for url in zip_links:
    print(f"Processing: {url}")

    response = requests.get(url, verify=False)
    zip_bytes = io.BytesIO(response.content)

    with zipfile.ZipFile(zip_bytes) as z:
        txt_files = [f for f in z.namelist() if f.endswith(".txt")]

```

```

if not txt_files:
    print("No txt file found inside ZIP:", url)
    continue

txt_filename = txt_files[0]
text_data = z.read(txt_filename).decode("latin-1")

# Parse tables
tables = parse_fama_french_text(text_data)

year_match = re.search(r"/.*?(\d{4}).*?/", url)
if year_match:
    year = year_match.group(1)
else:
    print("Could not extract year → skipping:", url)
    continue

year_folder = os.path.join(output_dir, year)
os.makedirs(year_folder, exist_ok=True)

# Save CSV files inside YEAR folder
for name, df in tables.items():
    clean_name = name.replace(" ", "_").replace("-", "")
    csv_path = os.path.join(year_folder, f"{clean_name}.csv")

    df.to_csv(csv_path, index=False)
    print(f"Saved → {csv_path}")

print("\nAll files processed successfully")

```

Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data\_Library/Historical\_Archives/08\_2017\_Update/ftp/6\_Portfolios\_2x3\_TXT.zip

C:\Users\nshinga\without\_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: <https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings>  
warnings.warn(

Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data\_Library/Historical\_Archives/08\_2018\_Update/ftp/6\_Portfolios\_2x3\_TXT.zip

```
C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2018\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2018\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2019 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2019\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2019\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2020 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2020\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2020\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2021 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2021\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2021\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2022 Update/ftp/6_Portfolios_2x3_TXT.zip
```

```
C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2022\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2022\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2023 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2023\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2023\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2024 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2024\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2024\Average_Equal_Weighted_Returns_Monthly.csv
Processing: https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/Data_Library/Historical_Archives/08 2025 Update/ftp/6_Portfolios_2x3_TXT.zip

C:\Users\nshinga\without_conda\Lib\site-packages\urllib3\connectionpool.py:1099: InsecureRequestWarning: Unverified HTTPS request is being made to host 'mba.tuck.dartmouth.edu'. Adding certificate verification is strongly advised. See: https://urllib3.readthedocs.io/en/latest/advanced-usage.html#tls-warnings
    warnings.warn(
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2025\Average_Value_Weighted_Returns_Monthly.csv
Saved → C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2025\Average_Equal_Weighted_Returns_Monthly.csv
```

All files processed successfully

**upload data to s3**

```
In [6]: local_base_dir = r"C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv"
bucket_name = "portfolio-return-data"
s3_prefix = "ff_6_portfolios/output_csv"

for root, dirs, files in os.walk(local_base_dir):
    for file in files:
        if file.endswith(".csv"):
            local_path = os.path.join(root, file)

            rel_path = os.path.relpath(local_path, local_base_dir)
            s3_key = os.path.join(s3_prefix, rel_path).replace("\\", "/")

            print(f"Uploading {local_path} → s3://{bucket_name}/{s3_key}")
            s3.Bucket('portfolio-return-data').upload_file(local_path, s3_key)

print("Upload complete")
```



```
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2025\Average_Equal_Weighted_Returns_Monthly.csv → s3://portfolio-return-data/ff_6_portfolios/output_csv/2025/Average_Equal_Weighted_Returns_Monthly.csv
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2025\Average_Value_Weighted_Returns_Monthly.csv → s3://portfolio-return-data/ff_6_portfolios/output_csv/2025/Average_Value_Weighted_Returns_Monthly.csv
Upload complete
```

## 2. Transform to Long Format

```
In [7]: main_path = r"C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv"

raw_data_list = []
ingestion_map_list = []

for root, dirs, files in os.walk(main_path):
    for file in files:
        if file.endswith('.csv'):
            file_path = os.path.join(root, file)
            df = pd.read_csv(file_path)

            # Convert to long format
            df = df.melt(id_vars=df.columns[0], var_name='PORTFOLIO', value_name='RETURN')

            # Convert date column to datetime
            df[df.columns[0]] = pd.to_datetime(df[df.columns[0]], format='%Y-%m-%d', errors='coerce')
            df.rename(columns={df.columns[0]: 'DATE'}, inplace=True)

            # Add metadata
            df['SOURCE'] = file_path
            df['INGESTION_TIME'] = datetime.now()
            df['INGESTED_BY'] = 'Neha Shingare'
            df['INGESTED_SYSTEM'] = 'Python_Local'

            raw_data_list.append(df)

            # Record ingestion map
            ingestion_map_list.append({
                'FILE_NAME': file,
                'YEAR_FOLDER': os.path.basename(root),
                'INGESTION_TIME': datetime.now(),
```

```
'RECORDS_INGESTED': len(df)
})
```

### 3. Access method and Source Map Table

```
In [8]: %%time
RAW_DATA = pd.concat(raw_data_list, ignore_index=True)
DATA_INGESTION_MAP = pd.DataFrame(ingestion_map_list)

# Cleaning data
RAW_DATA['DATE'] = pd.to_datetime(RAW_DATA['DATE'])
RAW_DATA['RETURN'] = pd.to_numeric(RAW_DATA['RETURN'], errors='coerce')

# handling nulls & duplicates
RAW_DATA.dropna(subset=['DATE', 'PORTFOLIO', 'RETURN'], inplace=True)
RAW_DATA.drop_duplicates(inplace=True)

pit_data_list = []
for year in range(2017, 2026):
    pit_df = RAW_DATA[RAW_DATA['DATE'] <= pd.Timestamp(year=year, month=7, day=31)].copy()
    pit_df['PIT_YEAR'] = year
    pit_data_list.append(pit_df)
PIT_DATA = pd.concat(pit_data_list, ignore_index=True)

# CURRENT_DATA = RAW_DATA.copy()
CURRENT_DATA = RAW_DATA.loc[RAW_DATA.groupby('PORTFOLIO')['DATE'].idxmax()].reset_index(drop=True)

output_path = r"C:\Users\nshinga\OneDrive - MORNİNGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output"
os.makedirs(output_path, exist_ok=True)

RAW_DATA.to_csv(os.path.join(output_path, 'RAW_DATA.csv'), index=False)
PIT_DATA.to_csv(os.path.join(output_path, 'PIT_DATA.csv'), index=False)
CURRENT_DATA.to_csv(os.path.join(output_path, 'CURRENT_DATA.csv'), index=False)
DATA_INGESTION_MAP.to_csv(os.path.join(output_path, 'DATA_INGESTION_MAP.csv'), index=False)

print("Data processing complete")
```

Data processing complete  
CPU times: total: 10.1 s  
Wall time: 27.6 s

```
In [9]: RAW_DATA.shape, PIT_DATA.shape, CURRENT_DATA.shape
```

```
Out[9]: ((110112, 7), (973728, 8), (6, 7))
```

## 4. Upload data to S3

```
In [10]: %%time
local_dir = r"C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output"
bucket_name = "portfolio-return-data"
s3_prefix = "ff_6_portfolios/processed_output/"

for file in os.listdir(local_dir):
    if file.endswith(".csv"):
        local_path = os.path.join(local_dir, file)
        s3_key = os.path.join(s3_prefix, file).replace("\\", "/")

        print(f"Uploading {local_path} → s3://{bucket_name}/{s3_key}")
        s3.Bucket('portfolio-return-data').upload_file(local_path, s3_key)

print("Upload complete")
```

```
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output\CURRENT_DATA.csv → s3://portfolio-return-data/ff_6_portfolios/processed_output/CURRENT_DATA.csv
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output\DATA_INGESTION_MAP.csv → s3://portfolio-return-data/ff_6_portfolios/processed_output/DATA_INGESTION_MAP.csv
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output\PIT_DATA.csv → s3://portfolio-return-data/ff_6_portfolios/processed_output/PIT_DATA.csv
Uploading C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\processed_output\RAW_DATA.csv → s3://portfolio-return-data/ff_6_portfolios/processed_output/RAW_DATA.csv
Upload complete
CPU times: total: 12.8 s
Wall time: 28.2 s
```

## Data quality check

```
In [21]: def robustness_checks(df_raw):
    print("1. NULL CHECKS")
    null_counts = df_raw.isnull().sum()
    print(null_counts)
```

```
print("\n")

print("2. DUPLICATE CHECKS")
dup_count = df_raw.duplicated().sum()
print(f"Total duplicated rows: {dup_count}")
if dup_count > 0:
    print("Removing duplicates...")
    df_raw.drop_duplicates(inplace=True)
print("\n")

print("3. SCHEMA & DTYPES VALIDATION")
expected_columns = ['DATE', 'PORTFOLIO', 'RETURN', 'SOURCE', 'INGESTION_TIME', 'INGESTED_BY', 'INGESTED_SYSTEM']

missing_cols = [col for col in expected_columns if col not in df_raw.columns]
if missing_cols:
    print(f"Missing required columns: {missing_cols}")
else:
    print("All required columns present")

# enforce types
try:
    df_raw['DATE'] = pd.to_datetime(df_raw['DATE'])
    print("DATE converted to datetime")
except:
    print("Failed to convert DATE to datetime")

try:
    df_raw['RETURN'] = pd.to_numeric(df_raw['RETURN'], errors='coerce')
    print("RETURN converted to numeric")
except:
    print("Failed to convert RETURN to numeric")

print("\n")

print("4. DATE RANGE VALIDATION")
min_date = df_raw['DATE'].min()
max_date = df_raw['DATE'].max()
print(f"Earliest DATE: {min_date}")
print(f"Latest DATE: {max_date}")

# Business rule: Data must be >= July 2017
cutoff_date = pd.Timestamp("2017-07-01")
```

```

if min_date < cutoff_date:
    print(f"Some dates earlier than July 2017 found: {min_date}")
else:
    print("No dates before July 2017")

# No future dates allowed
today = pd.Timestamp.today()
if max_date > today:
    print(f"Future dates detected: {max_date}")
else:
    print("No future dates detected")
print("\n")

print("5. MONTHLY CONTINUITY CHECK")
expected_months = pd.date_range(start=min_date, end=max_date, freq='MS')
actual_months = df_raw['DATE'].dt.to_period('M').unique()

missing_months = set(expected_months.to_period('M')) - set(actual_months)

if missing_months:
    print(f"Missing months: {sorted(missing_months)}")
else:
    print("No missing months in the dataset")
print("\n")

print("6. PORTFOLIO VALIDATION")
allowed_prefixes = ['SL', 'SM', 'SH', 'BL', 'BM', 'BH', 'Bi', 'Sm'] # Example for 6 FF portfolios

invalid_portfolios = df_raw[~df_raw['PORTFOLIO'].str[:2].isin([x[:2] for x in allowed_prefixes])]['PORTFOLIO'].unique()

if len(invalid_portfolios) > 0:
    print(f"Unexpected portfolio labels found: {invalid_portfolios}")
else:
    print("Portfolio labels follow expected pattern")
print("\n")

print("7. RETURN VALUE VALIDATION")
if df_raw['RETURN'].isnull().sum() > 0:
    print(f"▲ RETURN contains {df_raw['RETURN'].isnull().sum()} null values")

```

```
if not df_raw['RETURN'].between(-100, 100).all():
    print("Some RETURN values fall outside typical % return bounds (-100 to 100)")
else:
    print("RETURN values within reasonable range")
print("\n")

print("8. PIT (Point-In-Time) INTEGRITY CHECK")
# Example: ensure PIT cutoff logic holds:
# PIT for year X should only contain dates <= July of year X
try:
    PIT_DATA # check if exists
    for year in range(2017, 2026):
        pit_violations = df_raw[df_raw['DATE'] > pd.Timestamp(year=year, month=7, day=31)]

        if len(pit_violations) > 0:
            print("PIT dataset contains rows beyond PIT cutoff.")
        else:
            print("PIT dataset passes integrity checks")
except:
    print("PIT dataset not loaded in this environment")
print("\n")

print("9. S3 UPLOAD CHECK (Optional)")
s3 = boto3.client(
    service_name = 's3',
    region_name = 'ap-south-1',
    aws_access_key_id='AKIAQ-----HT5W5',
    aws_secret_access_key = 'upWVXX6----ucuZu72IH----78CeMwb---jxka'
)

bucket_name = 'portfolio-return-data'
s3_path_prefix = 'ff_6_portfolios/processed_output/'

response = s3.list_objects_v2(Bucket=bucket_name, Prefix=s3_path_prefix)

if 'Contents' in response:
    print(f"Files in '{s3_path_prefix}':")
    for obj in response['Contents']:
        print(obj['Key'])
```

```
else:  
    print(f"No files found in '{s3_path_prefix}' in bucket '{bucket_name}'.")  
  
return df_raw
```

```
In [22]: df_raw_checked = robustness_checks(CURRENT_DATA)
```

**1. NULL CHECKS**

```
DATE      0
PORTFOLIO 0
RETURN    0
SOURCE    0
INGESTION_TIME 0
INGESTED_BY 0
INGESTED_SYSTEM 0
dtype: int64
```

**2. DUPLICATE CHECKS**

Total duplicated rows: 0

**3. SCHEMA & DTYPES VALIDATION**

All required columns present  
DATE converted to datetime  
RETURN converted to numeric

**4. DATE RANGE VALIDATION**

Earliest DATE: 2025-07-01 00:00:00  
Latest DATE: 2025-07-01 00:00:00  
No dates before July 2017  
No future dates detected

**5. MONTHLY CONTINUITY CHECK**

No missing months in the dataset

**6. PORTFOLIO VALIDATION**

Portfolio labels follow expected pattern

**7. RETURN VALUE VALIDATION**

RETURN values within reasonable range

**8. PIT (Point-In-Time) INTEGRITY CHECK**

PIT dataset passes integrity checks

9. S3 UPLOAD CHECK (Optional)  
Files in 'ff\_6\_portfolios/processed\_output':  
ff\_6\_portfolios/processed\_output/CURRENT\_DATA.csv  
ff\_6\_portfolios/processed\_output/DATA\_INGESTION\_MAP.csv  
ff\_6\_portfolios/processed\_output/PIT\_DATA.csv  
ff\_6\_portfolios/processed\_output/RAW\_DATA.csv

In [23]: df\_raw\_checked = robustness\_checks(RAW\_DATA)

### 1. NULL CHECKS

```
DATE      0
PORTFOLIO 0
RETURN    0
SOURCE    0
INGESTION_TIME 0
INGESTED_BY 0
INGESTED_SYSTEM 0
dtype: int64
```

### 2. DUPLICATE CHECKS

Total duplicated rows: 0

### 3. SCHEMA & DTYPE VALIDATION

All required columns present  
DATE converted to datetime  
RETURN converted to numeric

### 4. DATE RANGE VALIDATION

Earliest DATE: 1926-07-01 00:00:00  
Latest DATE: 2025-07-01 00:00:00  
Some dates earlier than July 2017 found: 1926-07-01 00:00:00  
No future dates detected

### 5. MONTHLY CONTINUITY CHECK

No missing months in the dataset

### 6. PORTFOLIO VALIDATION

Portfolio labels follow expected pattern

### 7. RETURN VALUE VALIDATION

RETURN values within reasonable range

### 8. PIT (Point-In-Time) INTEGRITY CHECK

PIT dataset passes integrity checks

9. S3 UPLOAD CHECK (Optional)  
Files in 'ff\_6\_portfolios/processed\_output':  
ff\_6\_portfolios/processed\_output/CURRENT\_DATA.csv  
ff\_6\_portfolios/processed\_output/DATA\_INGESTION\_MAP.csv  
ff\_6\_portfolios/processed\_output/PIT\_DATA.csv  
ff\_6\_portfolios/processed\_output/RAW\_DATA.csv

In [24]: df\_raw\_checked = robustness\_checks(PIT\_DATA)

### 1. NULL CHECKS

```
DATE          0
PORTFOLIO    0
RETURN        0
SOURCE        0
INGESTION_TIME 0
INGESTED_BY   0
INGESTED_SYSTEM 0
PIT_YEAR      0
dtype: int64
```

### 2. DUPLICATE CHECKS

Total duplicated rows: 0

### 3. SCHEMA & DTYPES VALIDATION

All required columns present  
DATE converted to datetime  
RETURN converted to numeric

### 4. DATE RANGE VALIDATION

Earliest DATE: 1926-07-01 00:00:00  
Latest DATE: 2025-07-01 00:00:00  
Some dates earlier than July 2017 found: 1926-07-01 00:00:00  
No future dates detected

### 5. MONTHLY CONTINUITY CHECK

No missing months in the dataset

### 6. PORTFOLIO VALIDATION

Portfolio labels follow expected pattern

### 7. RETURN VALUE VALIDATION

RETURN values within reasonable range

### 8. PIT (Point-In-Time) INTEGRITY CHECK

PIT dataset passes integrity checks

9. S3 UPLOAD CHECK (Optional)  
Files in 'ff\_6\_portfolios/processed\_output/':  
ff\_6\_portfolios/processed\_output/CURRENT\_DATA.csv  
ff\_6\_portfolios/processed\_output/DATA\_INGESTION\_MAP.csv  
ff\_6\_portfolios/processed\_output/PIT\_DATA.csv  
ff\_6\_portfolios/processed\_output/RAW\_DATA.csv

```
In [25]: database_name = "ff_6_portfolios"
s3_data_location = "s3://portfolio-return-data/ff_6_portfolios/processed_output/"
s3_output = "s3://portfolio-return-data/athena-query-results/" # Athena query results

athena_client = boto3.client(
    service_name = 'athena',
    region_name = 'ap-south-1',
    aws_access_key_id='AKIAQ-----HT5W5',
    aws_secret_access_key = 'upWVXX6----ucuZu72IH----78CeMwb---jxka'
)

s3_client = boto3.client(
    service_name = 's3',
    region_name = 'ap-south-1',
    aws_access_key_id='AKIAQ-----HT5W5',
    aws_secret_access_key = 'upWVXX6----ucuZu72IH----78CeMwb---jxka'
)

def wait_for_query(query_execution_id):
    while True:
        response = athena_client.get_query_execution(QueryExecutionId=query_execution_id)
        status = response['QueryExecution']['Status']['State']
        if status in ["SUCCEEDED", "FAILED", "CANCELLED"]:
            break
        time.sleep(2)
    return status

# Create database
query = f"""
CREATE DATABASE IF NOT EXISTS {database_name}
LOCATION '{s3_data_location}'
```

```

"""
response = athena_client.start_query_execution(
    QueryString=query,
    ResultConfiguration={"OutputLocation": s3_output}
)
wait_for_query(response["QueryExecutionId"])
print(f"Database '{database_name}' ensured.")

# List CSV files in S3 folder
parsed_url = urllib.parse.urlparse(s3_data_location)
bucket = parsed_url.netloc
prefix = parsed_url.path.lstrip("/")

csv_objects = s3_client.list_objects_v2(Bucket=bucket, Prefix=prefix)
csv_files = [obj['Key'] for obj in csv_objects.get('Contents', []) if obj['Key'].endswith('.csv')]

if not csv_files:
    print("No CSV files found in the specified S3 folder.")
    exit()

# Create tables from CSVs
for csv_file in csv_files:
    table_name = csv_file.split("/")[-1].replace(".csv", "").replace("-", "_")

    obj = s3_client.get_object(Bucket=bucket, Key=csv_file)
    df = pd.read_csv(io.BytesIO(obj['Body'].read()), nrows=5)

    columns = ",\n        ".join([f"{col.replace(' ', '_')} string" for col in df.columns])

    # Location = parent folder, not the file itself
    s3_file_folder = s3_data_location # points to folder, not individual file

    create_table_query = f"""
CREATE EXTERNAL TABLE IF NOT EXISTS {database_name}.{table_name} (
    {columns}
)
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe'
WITH SERDEPROPERTIES (
    'serialization.format' = ',',
    'field.delim' = ','
)
"""

```

```
STORED AS TEXTFILE
LOCATION '{s3_file_folder}'
TBLPROPERTIES ('has_encrypted_data='false');
"""

response = athena_client.start_query_execution(
    QueryString=create_table_query,
    ResultConfiguration={"OutputLocation": s3_output}
)

status = wait_for_query(response["QueryExecutionId"])
if status == "SUCCEEDED":
    print(f"Table '{table_name}' created successfully.")
else:
    print(f"Failed to create table '{table_name}'. Status: {status}")
```

Database 'ff\_6\_portfolios' ensured.  
Table 'CURRENT\_DATA' created successfully.  
Table 'DATA\_INGESTION\_MAP' created successfully.  
Table 'PIT\_DATA' created successfully.  
Table 'RAW\_DATA' created successfully.

In [ ]:

In [ ]:

## AWS Services Used:

### 1. Amazon S3 — Data Lake Storage

Amazon S3 is used as the central storage layer for:

- **Raw data** (RAW\_DATA)
- **Point-in-Time snapshots** (PIT\_DATA)
- **Latest curated data** (CURRENT\_DATA)
- **Ingestion metadata** (DATA\_INGESTION\_MAP)

The screenshot shows the AWS S3 console interface. The left sidebar has a tree view under 'Amazon S3' with sections like 'General purpose buckets', 'Storage Lens', and 'Feature spotlight'. The main area shows a folder named 'processed\_output/' with four objects listed:

Name	Type	Last modified	Size	Storage class
CURRENT_DATA.csv	csv	November 25, 2025, 19:01:30 (UTC+05:30)	1.3 KB	Standard
DATA_INGESTION_MAP.csv	csv	November 25, 2025, 19:01:30 (UTC+05:30)	1.3 KB	Standard
PIT_DATA.csv	csv	November 25, 2025, 19:01:30 (UTC+05:30)	206.9 MB	Standard
RAW_DATA.csv	csv	November 25, 2025, 19:01:50 (UTC+05:30)	22.9 MB	Standard

### Reasons for choosing S3:

- Highly durable.
- Low cost for storing historical datasets.
- Serverless and fully managed.
- Integrates directly with Athena, Redshift Spectrum, Glue, and other AWS analytics services.
- Ideal for a *data lake* pattern where files are append-only and immutable.

### Assumptions / Constraints:

- Dataset size is small (few MBs).
- Monthly ingestion frequency.
- Low query concurrency.

## 2. AWS Glue Data Catalog — Metadata Layer

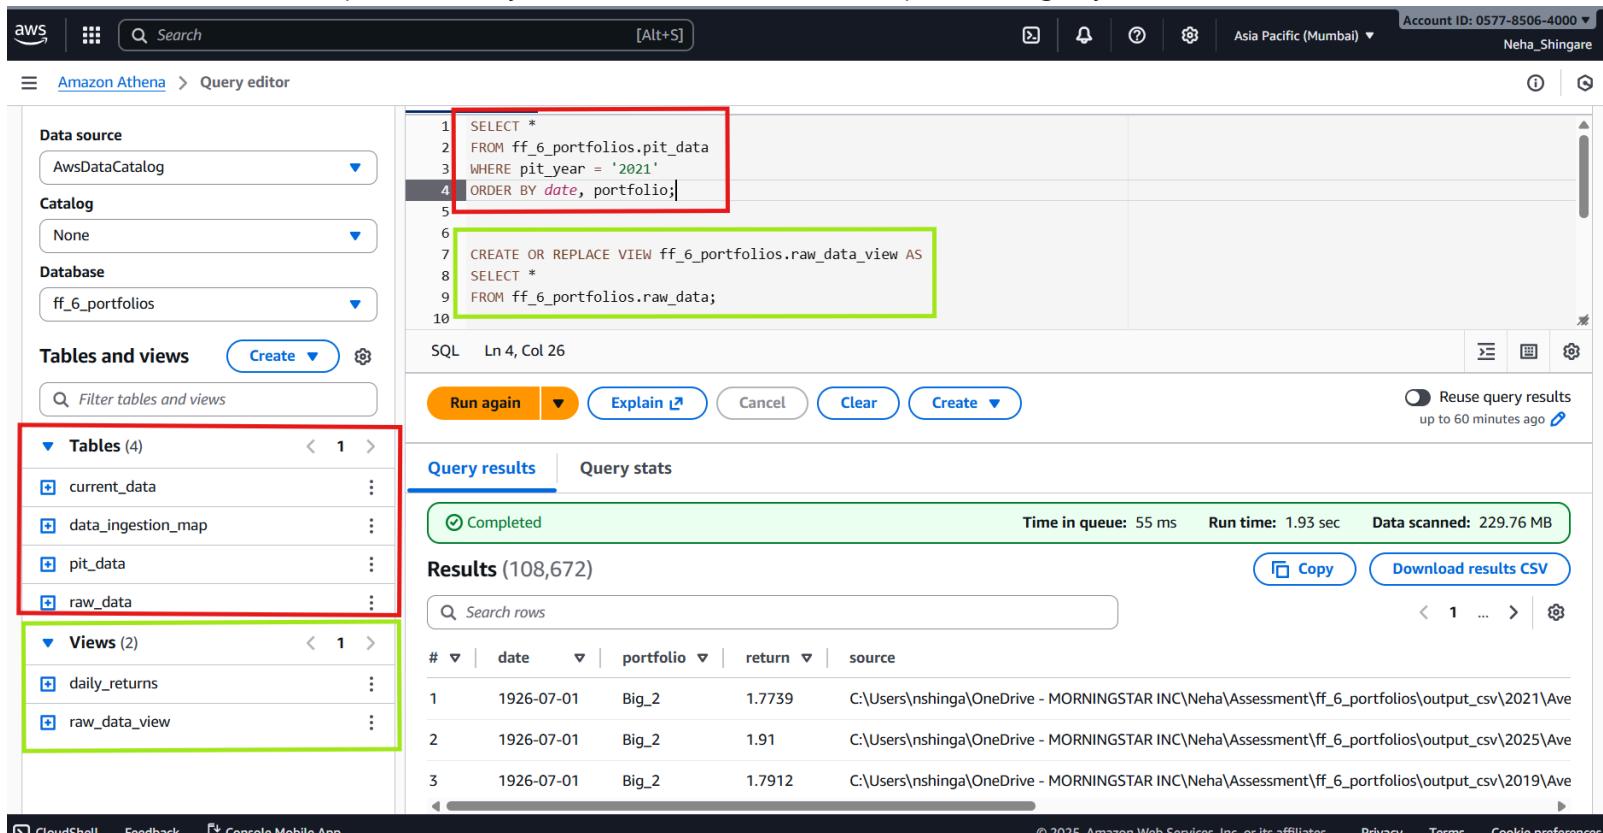
Athena uses the Glue Data Catalog to store table definitions for:

- RAW\_DATA
- PIT\_DATA
- CURRENT\_DATA
- DATA\_INGESTION\_MAP

This provides schema-on-read capability over files stored in S3.

## 3. Amazon Athena — Query Engine

Athena is used to run SQL queries directly on S3-based datasets without provisioning any servers.



The screenshot shows the Amazon Athena Query Editor interface. On the left, the sidebar displays the data source (AwsDataCatalog), catalog (None), and database (ff\_6\_portfolios). Below this, the 'Tables and views' section lists four tables: current\_data, data\_ingestion\_map, pit\_data, and raw\_data. It also lists two views: daily\_returns and raw\_data\_view. The main area contains a query editor with the following SQL code:

```

1 SELECT *
2 FROM ff_6_portfolios.pit_data
3 WHERE pit_year = '2021'
4 ORDER BY date, portfolio;
5
6
7 CREATE OR REPLACE VIEW ff_6_portfolios.raw_data_view AS
8 SELECT *
9 FROM ff_6_portfolios.raw_data;
10

```

The first four lines of the query are highlighted with a red box, and the last six lines are highlighted with a green box. Below the query editor, the 'Query results' tab is selected, showing the status 'Completed'. The results table has a header row with columns: #, date, portfolio, return, and source. The first three rows of the results table are as follows:

#	date	portfolio	return	source
1	1926-07-01	Big_2	1.7739	C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2021\Ave
2	1926-07-01	Big_2	1.91	C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2025\Ave
3	1926-07-01	Big_2	1.7912	C:\Users\nshinga\OneDrive - MORNINGSTAR INC\Neha\Assessment\ff_6_portfolios\output_csv\2019\Ave

--Tables and athena query in red box, view creation and views in green box

### Reasons for choosing Athena:

- Serverless: no cluster to manage.
- Pay-per-query: cost-effective for low/medium volumes.
- Perfect for analytical workloads over partitioned, columnar S3 data.
- Automatically scales with the dataset.

### Typical queries run:

- Validation of RAW\_DATA
- Querying PIT and CURRENT datasets
- Quality checks and auditing from the ingestion map

---

## Impact of Scaling Up (Task Takes 25 Minutes or Data Grows 100×)

If ingestion time grows significantly or dataset increases by two orders of magnitude:

### 1. Switch from CSV to Parquet

- 10–40× reduction in storage size
- Much faster Athena queries
- Lower S3 scan cost

### 2. Introduce S3 Partitioning

Partition RAW\_DATA by: year= / month=

Athena scans only relevant partitions → major performance improvement.

### 3. Move ETL to AWS Glue or AWS Lambda

- Glue ETL Jobs for large Spark-style transformations
- Lambda for lightweight parallel ingestion tasks

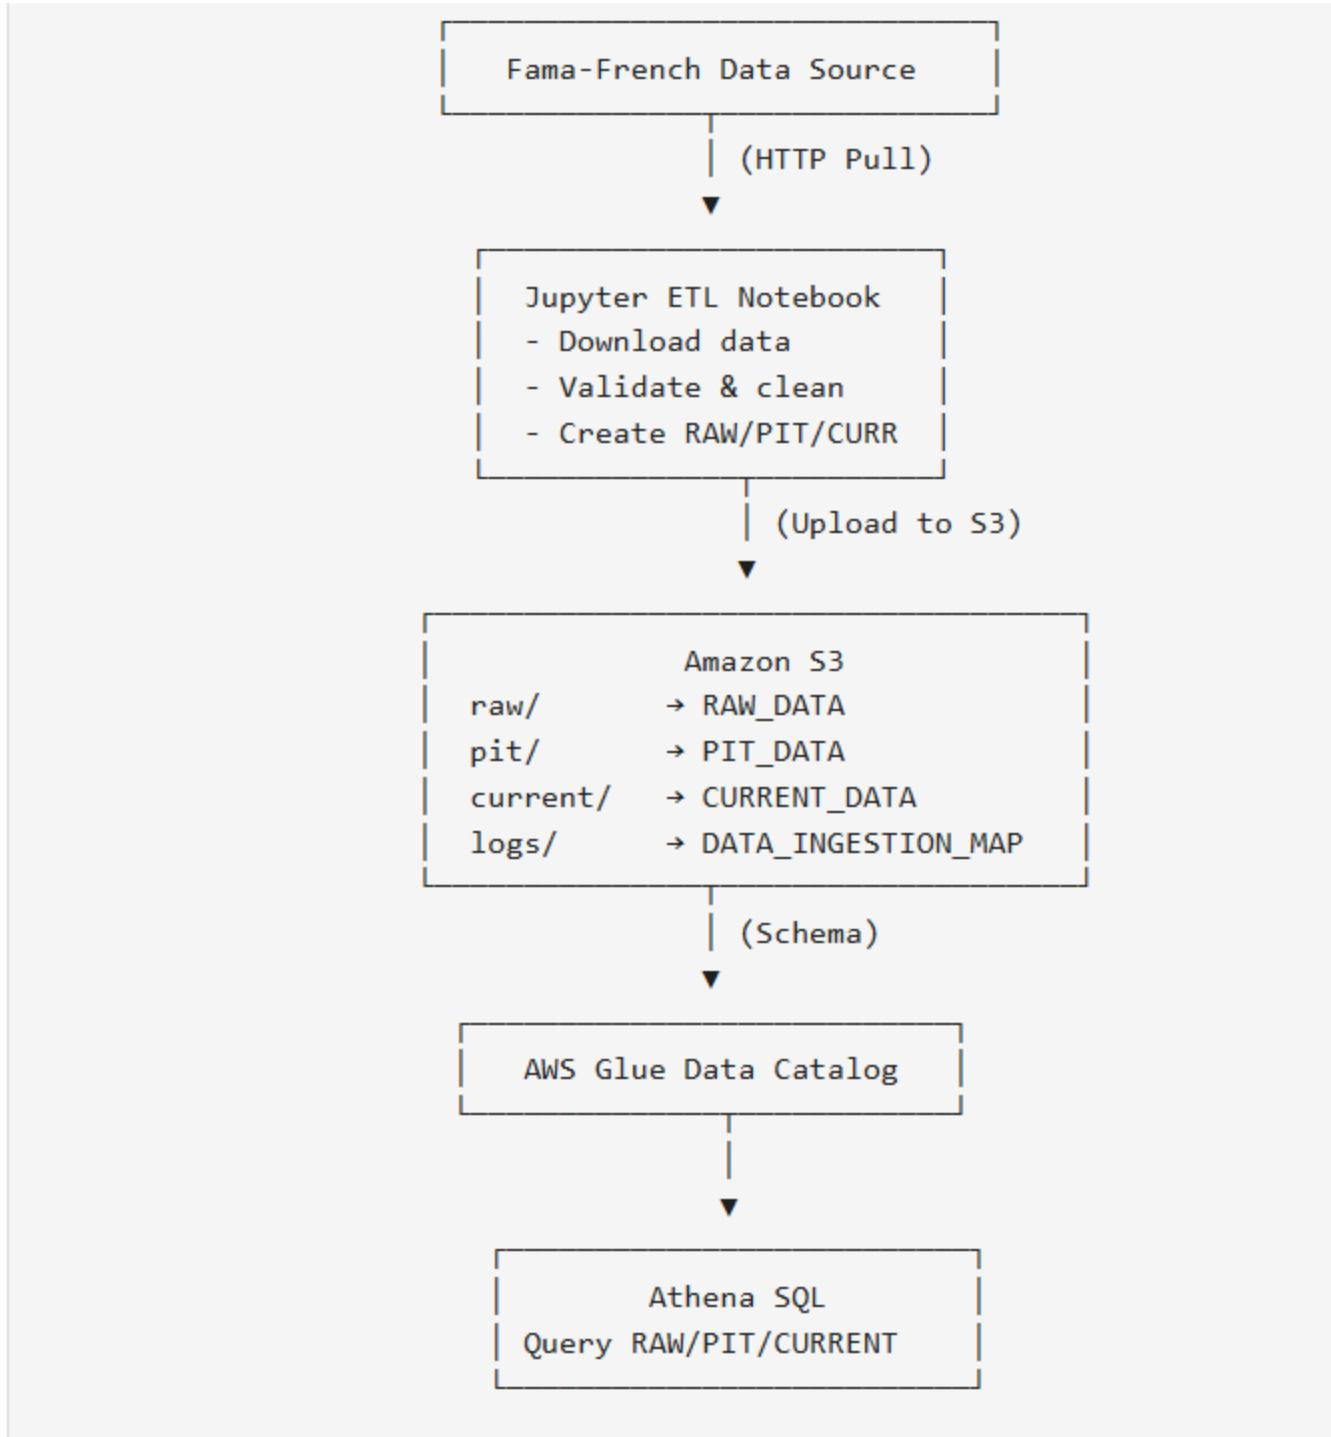
#### 4. Use Workflow Orchestration (AWS Step Functions)

To manage: Download → Transform → Validate → Upload → Update Metadata

#### 5. Consider Redshift for Heavy / Frequent Queries

If analytics workloads become intensive, migration to Redshift or using Redshift Spectrum is appropriate.

#### Architecture Diagram:



Optional for scale:

Amazon Redshift  
(via Redshift Spectrum)

## **Robustness Check:**

1. Null check
2. Duplicate check
3. Schema and dtype validation
4. Date range validation
5. Monthly Continuity Check
6. Portfolio Validation
7. Return Value Validation
8. S3 upload check

## **Notes:**

- ***Before running this notebook, please update the AWS configuration by inserting your own Access Key ID and Secret Access Key in the aws\_access\_key\_id and aws\_secret\_access\_key variables.***
- Naming conventions followed ( UPPER\_CASE for output tables, snake\_case for internal variables)
- Long format for scalability and easier pivoting
- Metadata for traceability ( ingestion\_time , source\_url , ingested\_by )
- PIT vs Current for reproducibility vs latest reporting

---

This notebook structure covers:

- Data ingestion
- Transformation to long format
- PIT & current data
- Source map / metadata

- Error handling & robustness
  - AWS services & architecture explanation
  - Outputs clearly defined
- 

In [ ]: